

Early Performance Analysis and Architectural Optimization of a Machine Learning Accelerator Using MatchLib

Michael Fingeroff | High-Level Synthesis Technologist
Siemens EDA, a part of Siemens Digital Industries Software

| Agenda

Challenges in Designing AI/ML Hardware
Convolutional Neural Network (CNN) Overview
Early Performance Analysis of CNN Convolution
Architectural Refinement
Synthesis and RTL Verification

AI/ML Application Challenges

Algorithmic Complexity

- Growing faster than the ability of RTL designers to code and verify

Memory Architecture Complexity



- Efficient data movement is key for power, performance and area

RTL Verification Costs Increasing

- Increased design complexity increases bugs introduced during hand-coding of RTL
- RTL regressions involve server farms, electricity cost, licenses and time

Slips in Design Schedule Kills Total Profit

- Finding bugs during system integration is too late



Convolutional Neural Network Overview

Convolutional Neural Network Overview

Used in object detection and classification

Mostly Convolutional layers

- Majority of computation done here (over 99%)
- Majority of memory traffic
- Bias and activation functions

Pooling layers

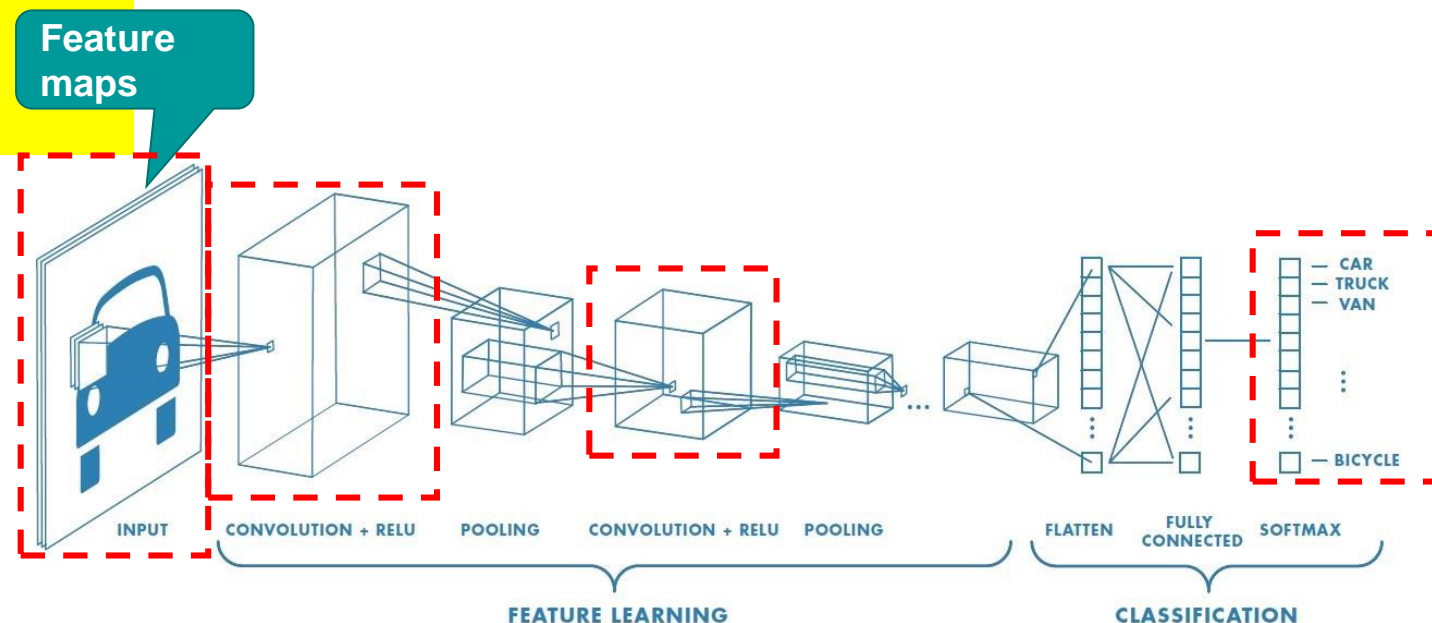
- Reduce feature map size

Fully connected

- Classification

Softmax

- normalize class probabilities



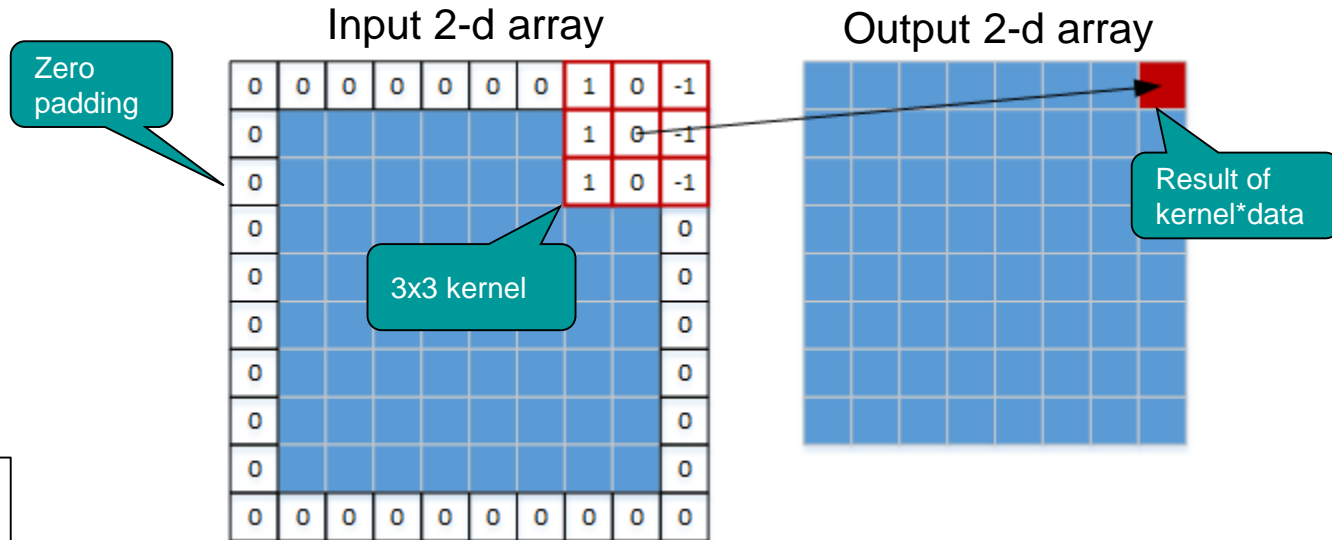
2-d 3x3 Convolution Algorithm

- 2-d kernel
- 2-d input array
- Padding left, right, and top and bottom
- Stride = number of elements kernel jumps by

2-d Convolution Algorithm – stride 1, zero pad

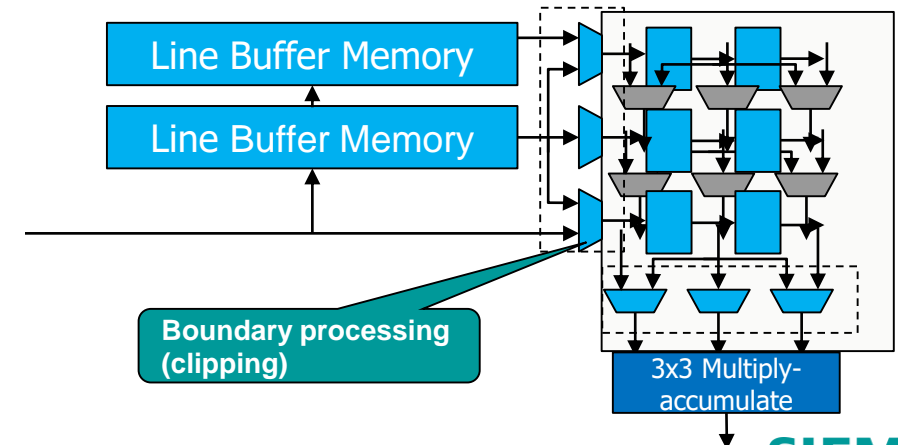
```

HEIGHT:for(int r=0;r<IN_HEIGHT;r++){
  WIDTH:for(int c=0;c<IN_WIDTH;c++){
    KERNEL_R:for(int i=0;i<3;i++){
      KERNEL_C:for(int j=0;j<3;j++){
        if(r+i-1 < 0 || r+i-1 >= IN_HEIGHT)
          data = 0;//zero pad
        else if(c+j-1 < 0 || c+j-1 >= IN_WIDTH)
          data = 0;//zero pad
        else
          data = data_in[r+i-1][c+j-1];
        acc += data * kernel[i][j];
      }
    }
    result[c] = acc;
  }
}
    
```



!=

Efficient Convolution Hardware Architecture



CNN Convolution – conv2d

- CNN convolutional layers have multiple input and output feature maps
- Each output feature map is a sum of separate convolutions across all input feature maps

Output feature map

Input feature map


2-d convolution

```
OUT_FMAP:for(int oc=0;oc<OUT_FMAPS;oc++){
  IN_FMAP:for(int ic=0;ic<IN_FMAPS;ic++){
    FMAP_HEIGHT:for(int r=0;r<IN_HEIGHT;r++){
      FMAP_WIDTH:for(int c=0;c<IN_WIDTH;c++){
        KERNEL_Y:for(int i=0;i<3;i++){
          KERNEL_X:for(int j=0;j<3;j++){
            acc[r][c] += fmap[ic][r-i/2][c-j/2] * kernel[ic][oc][i][j];
          }
        }
      }
      if(<last input fmap>)
        fmap_out[oc][r][c] = acc[r][c];
    }
  }
}
```

Lots of data movement for feature maps and weights

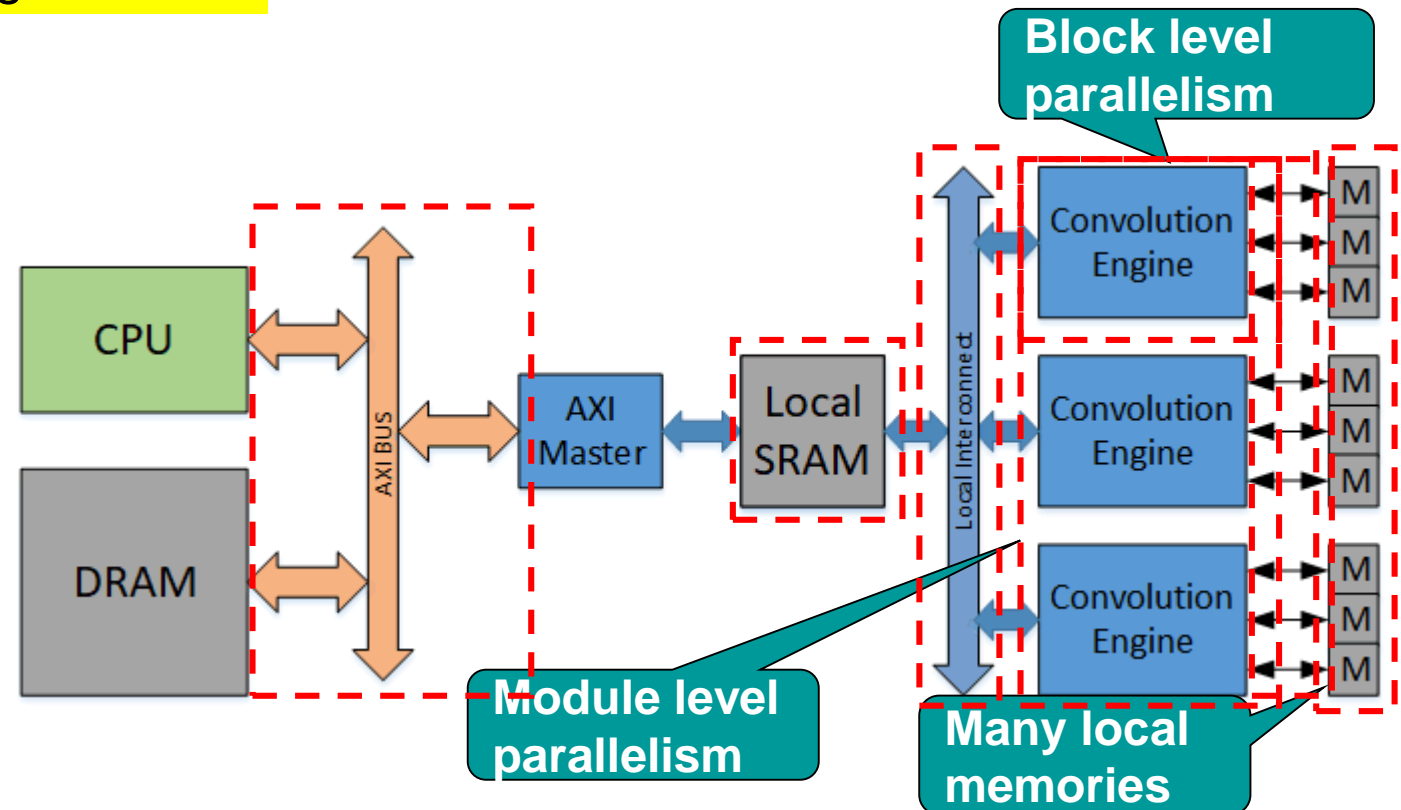
CNN Convolution – conv2d

- Every output fmap is 3-d convolution across the input fmaps
- Real CNNs have 100's or 1000's of fmaps
 - Lots of convolutions/computations
 - Lots of data movement



CNN Architectural Challenges

- Memory architectures need to leverage data reuse and parallelism
- May have multiple engines or processing elements
 - Block level parallelism
 - Module level parallelism
- Many local memories
- Complex interconnect
 - AXI4, local interconnect, etc.




| Coding Designs with MatchLib

Modelling Bus I/F With MatchLib

Users can easily model bus interfaces using MatchLib

- This example models a simple read bus I/F with burst




Initiate a burst by sending address and burst size

Read "burst_size" data from bus I/F

```
class dut : public sc_module {
public:
    ...
    sc_in<uint32>          addr_offset_csr;
    sc_in<uint8>          burst_csr;
    Connections::In <uint32> read_data_chan;
    Connections::Out <uint32> read_addr_chan;
    Connections::Out<uint32> read_burst_chan;
    ...
    void main() {
        wait();
        while (1) {
            uint32 addr = addr_offset_csr.read();
            uint8 burst_size = burst_csr.read();
            read_addr_chan.Push(addr);
            read_burst_chan.Push(burst_size);
            do {
                uint32 data = read_data_chan.Pop();
                read_data_out.Push(data);
            } while (--burst_size !=0);
        }
    }
}
```

Modelling Bus I/F With MatchLib

Small burst size and/or non-consecutive addresses will hurt performance by injecting dead cycles

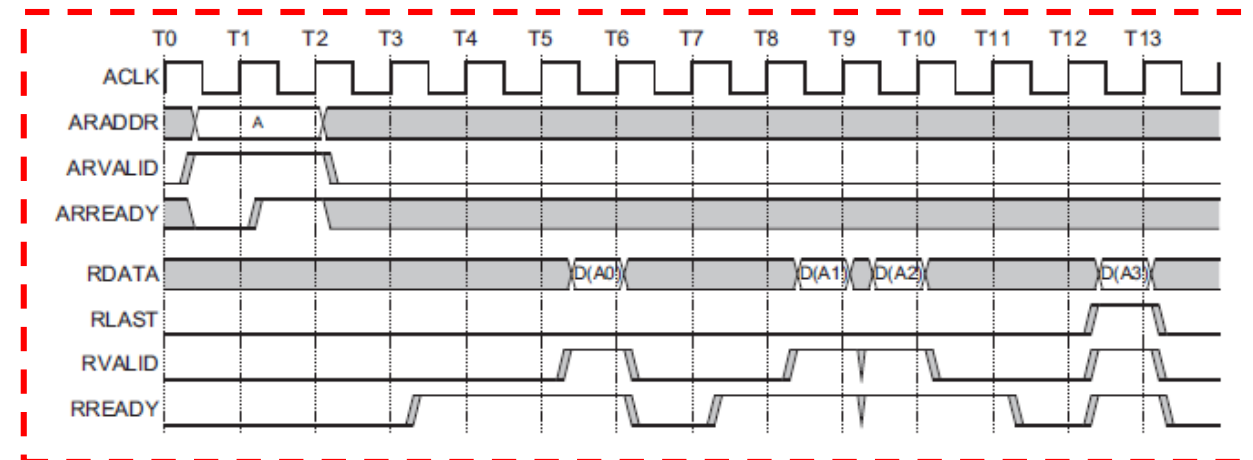


Complex Bus Protocols are Easily Modelled with MatchLib

- AXI4 Master and Slaves are part of the MatchLib library
 - AXI4 segmenter transparently manages burst size, 4k boundaries, etc
- AXI4 Read Master I/F with built-in segmentation
 - r_master<>
- AXI4 Write Master I/F with built-in segmentation
 - w_master<>
- AXI4 Slave I/F with built-in segmentation
 - slave<>
- Interface reads/writes use connections Push/Pop methods
 - Generates AXI4 transactions

```
class bus_if: public sc_module
...
    r_master<> r_master0;
...
    while (1) {
        ...
        r_payload r = r_master0.r.Pop();
        ...
    }
...

```





Early Performance Analysis of CNN Convolution

Design Goals

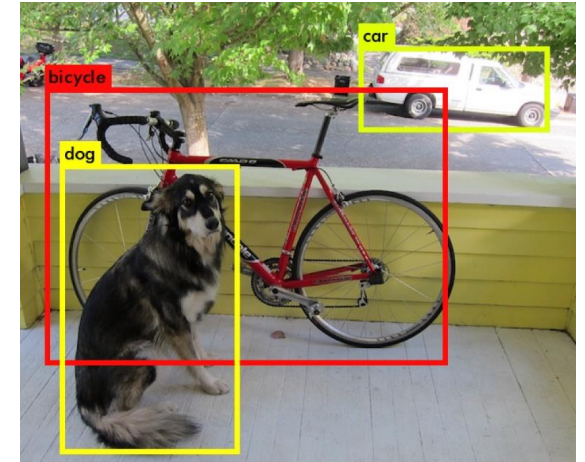
Implement a CNN for object detection and classification

- 9 layers
- Mostly 3x3 convolution (9 multiply-acc)
- 3.5 billion macs/inference

Low power/performance Ring-doorbell type application

- 1 inference/sec

500MHz clock frequency



Step 1: Original Algorithmic Model of conv2d


- Direct conversion of algorithm to HLS synthesizable bit-accurate model
- Generic bus interfaces with burst
 - Read burst size limited to one due to non-sequential addressing
 - Writes of feature maps can sustain large burst size
- No opportunity for parallelism

```
OFM:for(int ofm=0; ofm<OUT_FMAP; ofm++) {
  IFM:for (int ifm=0; ifm<IN_FMAP; ifm++) {
    ROW:for (int r=0; r<MAX_HEIGHT; r++) {
      COL:for (int c=0; c<MAX_WIDTH; c++) {
        K_X:for (int kr=0; kr<KSIZE; kr++) {
          K_Y:for (int kc=0; kc<KSIZE; kc++) {
            int ridx = r + kr - KSIZE/2;
            int cidx = c + kc - KSIZE/2;
            <zero pad>
            data_idx=rdoffset+ifm*ht*wt+ridx*wt+cidx;
            mem_in_addr.Push(data_idx);
            mem_in_burst.Push(1);
            data = mem_in_data.Pop();
            <weight read bus transaction>
            acc += data*mem_in_data.Pop();
          }
        }
      }
      acc_buf[r][c] += acc; ...
    }
  }
}
<Copy feature maps to system memory>
```

Step 1: Algorithmic Model Pre-HLS Simulation Results

- SystemC simulation “wall-clock time” took very long (~ 2 hours)
 - Context switching due to non-sequential memory accesses
 - Redundant memory accesses
- Simulation time was ~14 seconds to simulate 1 inference
- Poor design
 - No need to go any further

Pre-HLS Simulation

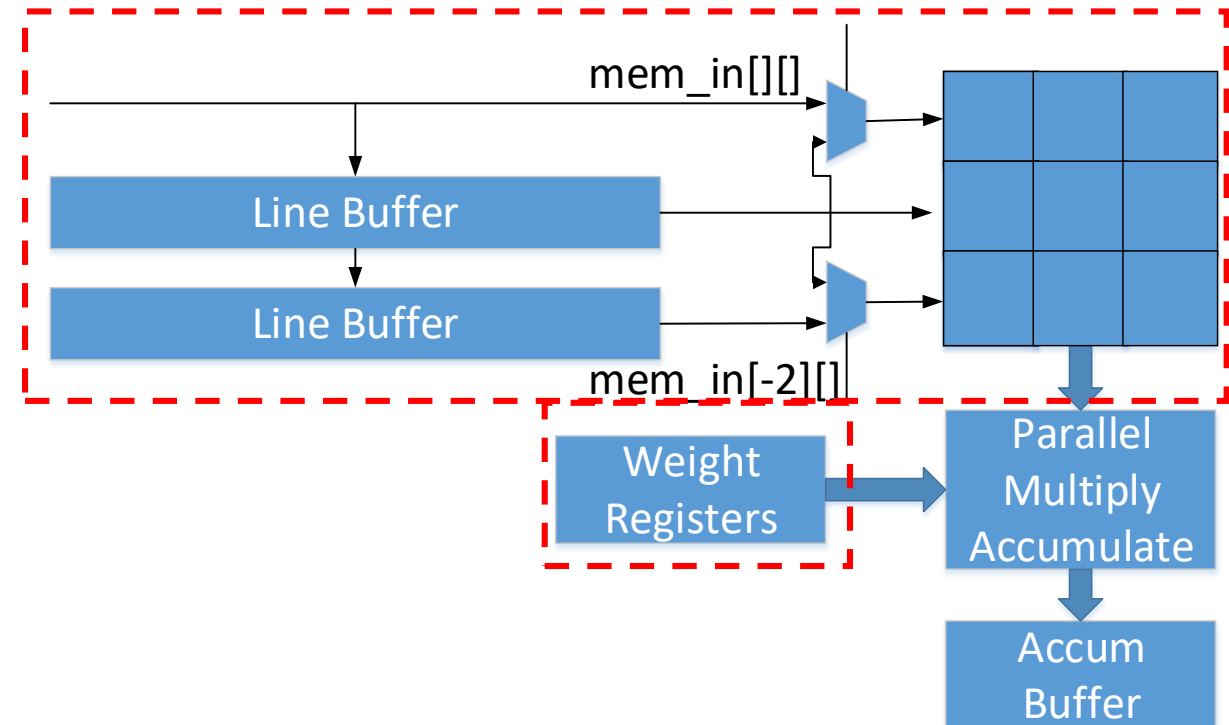




Architectural Refinement

Step 2: On-chip Buffering and Windowing

- SystemC HLS designs must be architected for efficient data movement and reuse
 - Improved simulation performance
 - Will allow HLS to extract parallelism
- Sliding-window architecture allows feature map data reuse
- Weight register cache read once for each input/output feature map computation



Step 2: On-chip Buffering and Windowing

- 9 weight bursts
 - Stored in register cache
- Feature maps burst a row at a time
 - Could also burst entire feature map
- Sliding window architecture allows K_X and K_Y to execute in parallel

```
OFM:for(int ofm=0;ofm<OUT_FMAP;ofm++){
  IFM:for(int ifm=0;ifm<IN_FMAP;ifm++){
    mem_in_addr.Push(weight_idx);
    mem_in_burst.Push(9);
    <cache weights>
    ROW:for(int r=0;r<MAX_HEIGHT+1;r++){
      data_idx=read_offset+ifm*height*width+r*width;
      if(r != height){
        mem_in_addr.Push(data_idx);
        mem_in_burst.Push(width);
      }
      COL:for(int c=0;c<MAX_WIDTH+1;c++){
        if(r != height && c != width)
          data[0] = mem_in_data.Pop();
        <sliding window architecture>
        K_X:for(int kr=0;kr<KSIZE;kr++){
          K_Y:for(int kc=0;kc<KSIZE;kc++){
            acc += window[kr][kc]*weights[kr][kc];
          }
        }
      }
    }
  }
}
```

“window” and
“weights” can be
read in parallel

Step 2: On-chip Buffering and Windowing Results

- Simulation results
 - Design goal met with simulation time of 0.864 secs
 - Pre-hls simulation wall-clock time 34 minutes for 1 inference
- All other operations run in software
 - Bias, RELU, max pooling, etc.
 - SystemC testbench runs in zero time
- What can MatchLib and SystemC tell us about the system-level performance?


CPU Software Function Calls

```
preprocessing()  
setup_layer_parameters()  
start_conv2d()  
<HW executing>  
wait_for_done()  
bias_add();  
leakyRelu()  
max_pooling()  
post_processing()
```

Step 3: Interaction with the CPU

- Target hardware platform
 - System memory is shared between the CPU and the conv2d accelerator
 - There is no CPU cache
- SystemC testbench models arbitrated memory between CPU and ML accelerator
 - Approximated CPU instruction execution and memory access time
- The performance of the accelerator is throttled by the CPU
 - Simulation time 2.6 secs
 - Simulation wall-clock time 63 minutes
 - Time spent converting from fixed-point to float

SystemC Simulation of One Inference



Step 4: Fusing Computational Layers

- Move Bias, ReLU, and max pooling into the accelerator
 - Cost little more in hardware area
- Can be coded into the design where feature map data is copied back to system memory
- Design simulation time 0.9 secs for one inference
- Pre-HLS simulation wall-clock time 30 minutes

```
<Get bias from system memory>
ROW_CPY:for (int r=0; r<MAX_HEIGHT+1; r++) {
    <setup burst size>
    mem_out_addr.Push(out_idx);
    mem_out_burst.Push(burst_size);
    COL_CPY:for (int c=0; c<MAX_WIDTH+1; c++) {
        add_bias = acc_buf[r][c] + bias;
        if (relu)
            if (add_bias < 0)
                add_bias = add_bias * SAT_TYPE(0.1);
        ...
        if(pool){
            <max_pooling>
            mem_out_data.Push(max);
        }else
            mem_out_data.Push(add_bias);
    }
}
```

Step4 – Bus Utilization

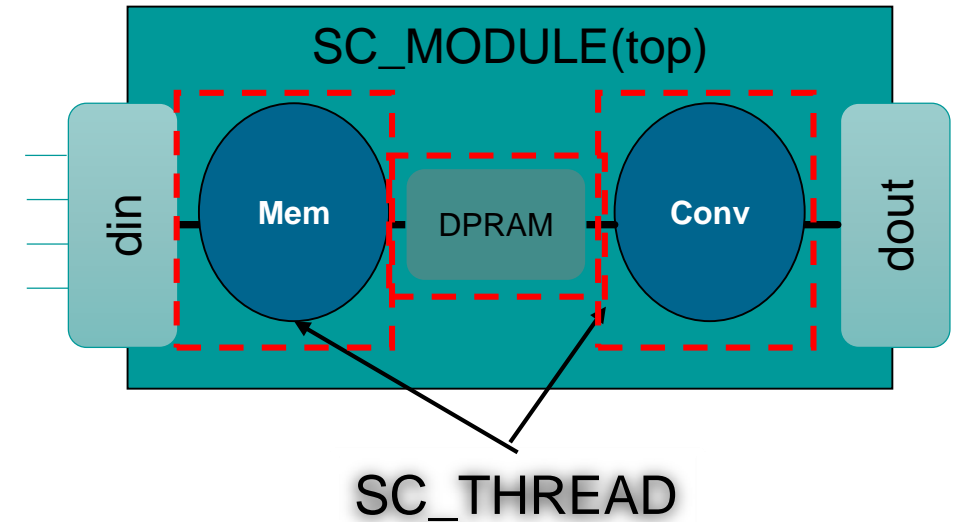
- Memory bus is ~100% utilized by the ML accelerator
- Input feature maps are re-read from system memory for each output feature map computation
 - System memory accesses are an order of magnitude larger for power consumption compared to on-chip SRAM

Pre-HLS Simulation



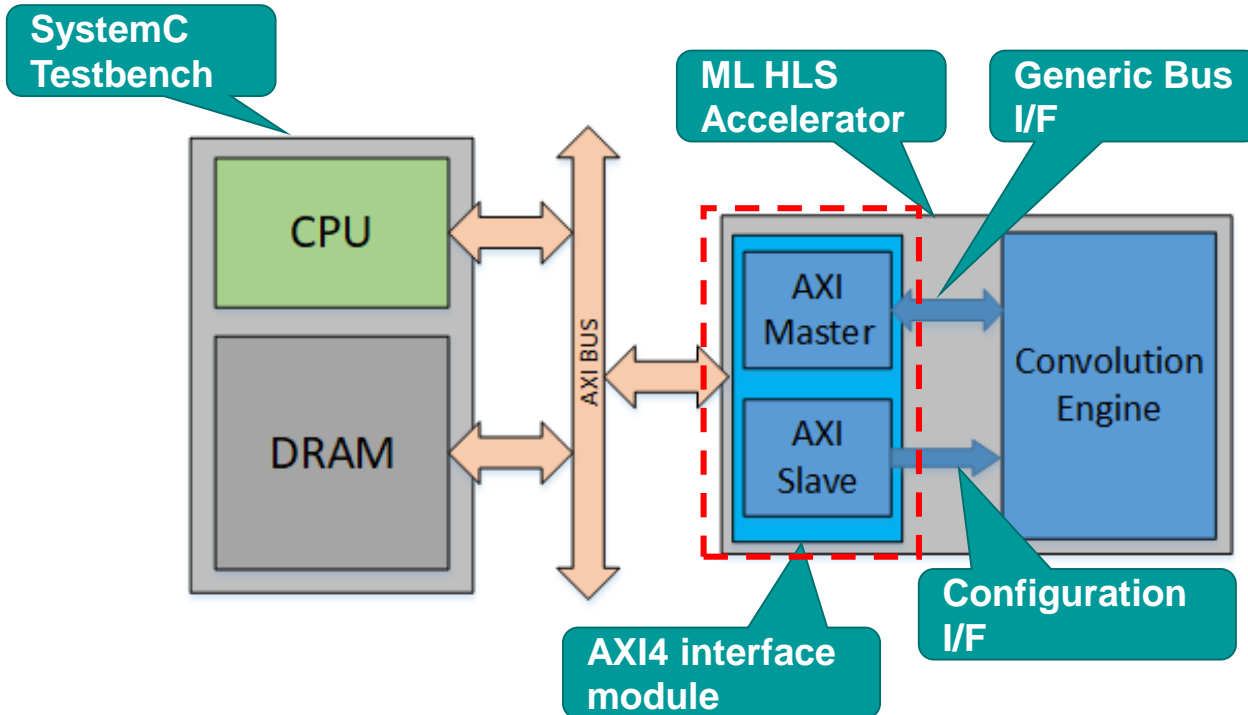
Step 5: Adding On-chip Buffering

- Buffer feature maps on-chip
 - ~800 KB for full buffering
 - SystemC memory generated by Catapult memory generator
- Split design into multiple processes
 - Memory read process to access system memory
 - Convolution, bias, ReLu, and max pooling process
 - Shared instantiated SystemC feature map memory between processes



Step 5: Adding AXI Master and Slave interfaces

- Existing design with generic bus I/F can be easily “bolted” to an AXI4 master and slave I/F module
- Generic I/F configures the AXI4 master I/F
 - Push/pop



```
class bus_interface : public sc_module, public local_axi {
public:
    sc_in<bool> clk);
    sc_in<bool> rstn);

    //AXI4 Bus I/Fs
    r_master<> r_master0;
    w_master<> w_master0;
    typename local_axi4_lite::write::template slave<> w_slave0;

    //User I/F to DUT
    //Read I/F
    Connections::In<MEM_ADDR_TYPE> mem_in_addr;
    Connections::In<BURST_TYPE> mem_in_burst;
    Connections::Out<DTYPE> mem_in_data;

    ...
    void read_master_process() {
    ...
        ar.ex_len = mem_in_burst.Pop();
        ar.addr = uint32(mem_in_addr.Pop())<<1;
        r_segment0_ex_ar_chan.Push(ar);
        while (1) {
            r_payload r = r_master0.r.Pop();
            data. = r.data;
            mem_in_data.Push(data);
            if (ar.ex_len-- == 0) { break; }
        }
    }
```


Step 5: Adding AXI4 I/Fs and On-chip Buffering

- Bus traffic dramatically reduced to reading input feature maps once
- Simulation time was 0.93 secs for one inference
- Simulation wall-clock time was 20 minutes

Pre-HLS Simulation





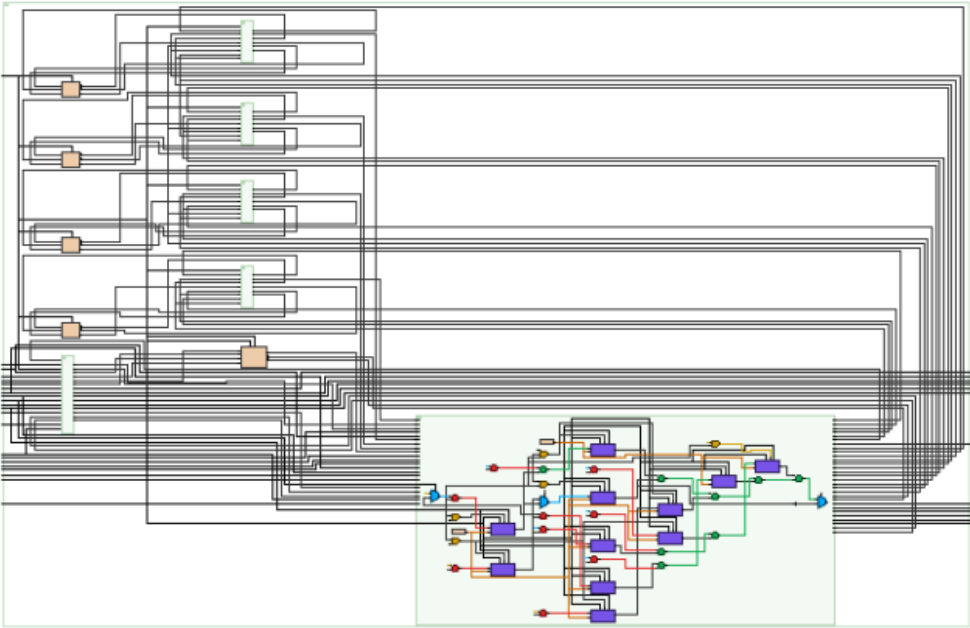
Synthesis and RTL Verification

Post-HLS Synthesis RTL Simulation Results

Post-HLS simulation results were very close to the pre-HLS simulation

Post-HLS simulation runtime was over 30x longer than the pre-HLS simulation

- Not practical for simulating multiple frames of video



Simulation Type	Simulation Time (secs)	Simulation Wall-clock time(mins)
Pre-HLS	0.93	20
Post-HLS	0.97	630

Source Code Steps

Source code examples and other tutorials can be found at:

- <https://hlslibs.org/>
- <https://github.com/hlslibs>

Summary

- Increasing AI/ML algorithm complexity is making RTL verification more difficult
- MatchLib and SystemC allows designers to model and verify the true hardware performance, catching bugs early that would normally be exposed during system integration when it's too late
- MatchLib models can be directly synthesized to RTL and performance of the pre-hls and post-hls results are near identical
- Customers are using MatchLib today to solve the design challenges associated with building AI/ML hardware

| Thank you!